

```
In [25]: # Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [26]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [27]: # Loading the dataset
data=pd.read_csv("C:/Users/HP/Desktop/project-DiabetesPrediction/1.diabetes_Data.csv")
data
```

```
Out[27]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	33
3	1	89	66	23	94	28.1	0.167	41
4	0	137	40	35	168	43.1	2.288	33
...
763	10	101	76	48	180	32.9	0.171	41
764	2	122	70	27	0	36.8	0.340	31
765	5	121	72	23	112	26.2	0.245	33
766	1	126	60	0	0	30.1	0.349	31
767	1	93	70	31	0	30.4	0.315	33

768 rows × 9 columns




```
In [28]: # Checking for missing values
data.isnull().sum()
```

```
Out[28]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
In [29]: #Checking for missing values  
sns.heatmap(data.isnull())
```

Out[29]: <Axes: >



```
In [12]: # Co relation matrix
correlation=data.corr()
print(correlation)
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	
Outcome	0.221898	0.466581	0.065068	0.074752	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
In [30]: # Visualising the Co relation  
sns.heatmap(data.corr())
```

Out[30]: <Axes: >



```
In [31]: # Train test split  
x=data.drop("Outcome",axis=1)  
y=data["Outcome"]  
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=21)
```

```
In [32]: # Training the model
         model=LogisticRegression()
         model.fit(x_train,y_train)
```

```
D:\Users\HP\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py:460: Co
nvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
(https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression)
n_iter_i = _check_optimize_result(
```

```
Out[32]: LogisticRegression
LogisticRegression()
```

```
In [33]: # Making Prediction
prediction = model.predict(x_test)
prediction
```

```
Out[33]: array([0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0,
                1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
                0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
                0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
                0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
                0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1,
                1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0,
                0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1,
                1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [34]: # Accuracy_score of the model
accuracy = accuracy_score(prediction,y_test)
accuracy
```

Out[34]: 0.7272727272727273

```
In [35]: from sklearn.metrics import confusion_matrix, classification_report
```

```
In [36]: #Confusion matrix
cm = confusion_matrix(y_test,prediction)
cm
```

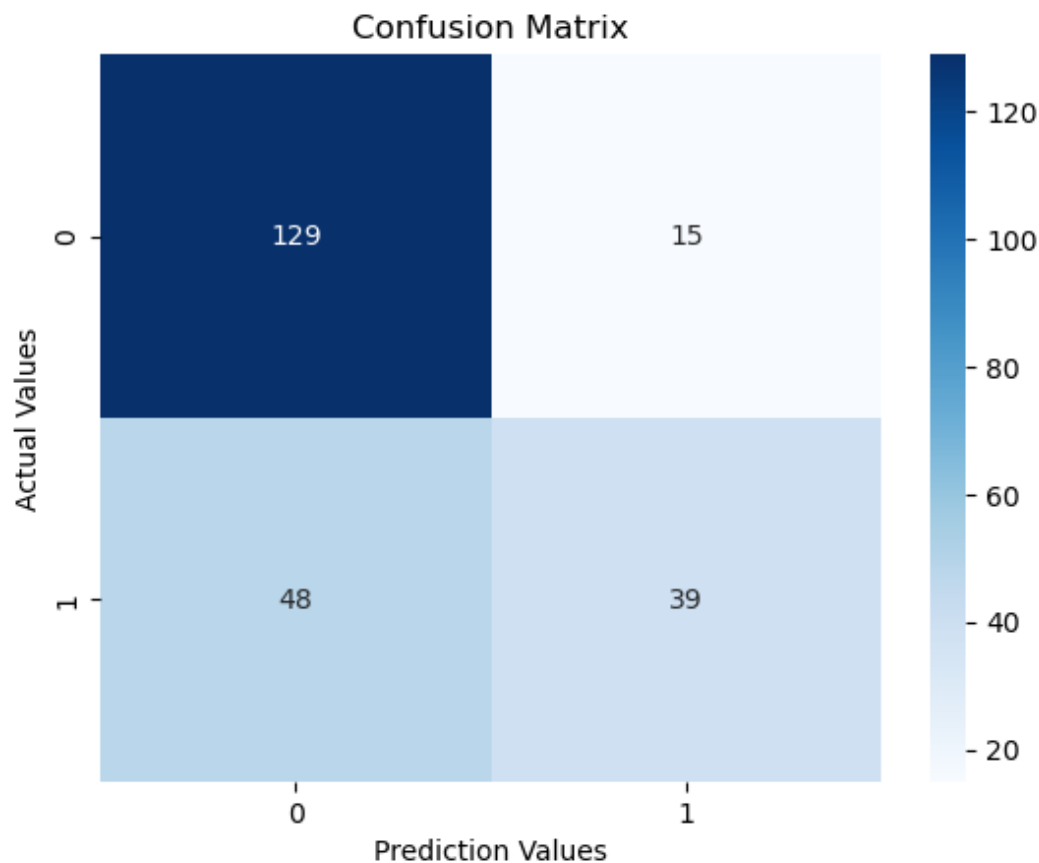
```
Out[36]: array([[129, 15],
                [ 48, 39]], dtype=int64)
```

```
In [37]: # Classification Report
print(classification_report(prediction,y_test))
```

	precision	recall	f1-score	support
0	0.90	0.73	0.80	177
1	0.45	0.72	0.55	54
accuracy			0.73	231
macro avg	0.67	0.73	0.68	231
weighted avg	0.79	0.73	0.75	231

```
In [40]: # Graphical Presentation of Confusion Matrix
plt.title("Confusion Matrix")
sns.heatmap(cm,annot = True,fmt='d',cmap = 'Blues')
plt.ylabel("Actual Values")
plt.xlabel("Prediction Values")
```

Out[40]: Text(0.5, 23.52222222222222, 'Prediction Values')



In []:

